PH244
Big Data: A Public Health Perspective
Computing Project

The file `ss13hus.csv.bz2` under `bCourses/Files/Project/Project-II/` contains household-specific data from the 2009-2013 US Census American Community Survey. This survey obtains a wealth of information on people and households every year, with about 1% of the total population surveyed in each year. The dictionary describing all the data fields is available as `PUMS-Data-Dictionary-2009-2013.pdf` under the same directory. The zipped file is about 600MB, and be careful about unzipping it. **You are required to use `R` for this computing project, and use `R Markdown` that include both the computer code and the output for the project report. There is no page limit on this report.**

1. Use `read.csv()` to read the zipped data `ss13hus.csv.bz2` into `R`. Considering the data size, use the "divide-and-conquer" strategy to read the data. To do so, first count the number of rows of the data using the function `countLines()`, then read a chunk of data each time, until reaching the last row of the data. Use `system.time()` to record and report the time it requires to read the data. The following code might be helpful:

   ```
   library(R.utils)
   nrow=countLines("ss13hus.csv.bz2")[1]
   filename <- "ss13hus.csv.bz2"
   con <- bzfile(filename, "rt")
   # choose a suitable value for chunk_size based on nrow
   DT1<-read.csv(con, header = T, nrows = chunk_size)
   ...
   ```

2. Create a subset of data by *randomly* sampling $3,000,000$ survey records from `ss13hus.csv.bz2`. Extract the following data fields: `REGION, ST, ADJHSG, ADJINC, NP, ACR, BDSP, ELEP, GASP, RMSP, VEH, WATP, FINCP, HINCP`. Save the file as a `csv` for subsequent analyses, with rows representing survey records and columns different data fields. In addition, for reproducibility, please use `set.seed(1000)` to set the random seed.

3. Try 3 different functions of reading the data you create in Step 2 into `R`: `read.csv()`, `fread()`, and `read.csv.ffdf()`. Use `system.time()` to record and report the time each function requires to read in the data.

4. Draw a scatter plot of `BDSP` (the number of bedrooms; a measure of house size) on the x-axis, and `FINCP` (the family income; use `ADJINC` to adjust `FINCP` to constant dollars) on the y-axis. Add a `gam` smoother, with standard error shading, on the scatter plot using the `R` package `ggplot2`.

5. Create a subset of data by *randomly* sample $1,000,000$ survey records from `ss13hus.csv.bz2` as a subset. Fit a linear regression model on this subset with the adjusted family income (`FINCP`) as the response, and `BDSP` and `VEH` (the number of vehicles) as the predictors. Record the estimated coefficient for `BDSP`. Repeat the same procedures for 1000 times with different random seeds. Use the `R` function `lm()` and `biglm()` to implement it. Use `system.time()` to report the time it requires in total. Report the mean and standard deviation of the estimated coefficients for `BDSP` recorded in these 1000 repetitions. Also draw a density plot of the estimated coefficients for `BDSP` recorded in these 1000 repetitions.